

claude-jacked-kids — Coverage Matrix Analysis

v2 · re-anchored to the Claude Code substrate

2026-06-28

Scope: main + PR #11 + shipped audio

A Claude Code plugin (skills + prompts + persona + README)

Frame correction (why this is v2): v1 scored this like a standalone kids' platform and docked it for "needs a terminal + a \$20/mo plan vs Scratch's web instant-on." Wrong denominator. This is a **Claude Code plugin** — its job is to help a kid *build using Claude Code itself*. Claude Code is the **given substrate** (like Windows is for a Windows app). The right bar is "best kid experience *on* the substrate," and the real competitor is **raw Claude Code with no plugin**. Re-anchoring lifts overall **4.56 → 5.17** and points the roadmap at the *real* gaps.

Coverage (re-anchored)

5.2

was 4.6
overall / 10

5.6

was 5.0
in-market

6.4

was 5.9
sweet spot
(Mid Kid)



Target outcomes

Outcome	Now	Target
Time-to-first-creation	?	<10 min
First-creation unassisted	?	≥80%
Repeat creations / kid / week	?	≥3


Unmeasured — the product is deliberately zero-telemetry. *That gap is itself a finding.*

🎧 **Shipped this session: an audio onboarding companion** — warm OpenAI-TTS "Listen instead" guides for a kid (who can't read the README yet) and a grown-up (60s setup), wired into the README + a persona backstop. Static `.mp3` s generated once; the kid's machine never calls a voice service (offline-first intact). This was the #3-ranked lever and directly lifts the weakest in-market row.

Contents

1. Product pipeline
2. The heatmap matrix
3. Key observations
4. Workflow experience
5. Competitive position
6. Cross-cutting levers
7. Phased roadmap
8. Target matrix
9. Strategic recommendations
10. Profiles
11. Sources

Product pipeline



Re-anchored: Setup is no longer the cold gate (the substrate is a given); Manage warms up (Claude Code's built-in `/rewind +git`). The remaining cool stages are Play & Show and Publish — the "show a friend" motivation loop.

The heatmap matrix

2 barely 3 basic 4 gaps 5 works, fights it 6 notable gaps 7 good 8 strong 9 excellent 10 best-in-class

Role \ Domain	Setup	Game	App	Change	Play	Learn	Identity	Publish	Manage	Avg
Young Kid (5–7)	4	4	3	3	4	4	4	3	4	3.7
Mid Kid (8–10)	6	8	7	7	6	8	6	4	6	6.4
Tween (11–12)	7	7	7	7	5	7	5	5	6	6.2
Teen (13–14)	7	6	6	6	5	5	4	6	6	5.7
Parent / Grown-up	7	6	6	6	6	7	6	7	5	6.2
Educator / Classroom	2	3	3	3	3	4	2	2	3	2.8
Avg	5.5	5.7	5.3	5.3	4.8	5.8	4.5	4.5	5.0	5.2

vs v1: Setup col **2.3→5.5** (substrate is a given), Manage **3.2→5.0** (CC's built-in `/rewind +git`), Publish **4.0→4.5**. Educator stays cold — classroom blockers (per-seat plans, no rostering/DPA) are real regardless of substrate. (Mid-Kid Setup/Publish scored to experienced value — 6/4 — per the §workflow table.)

Key observations

The asset — and it's first-of-its-kind. Build+teach for the sweet-spot kid (Make-a-Game 8, Learn 8). Research found **no kid-safe Claude Code plugin exists**, and *no* agent on the substrate teaches CS while it builds. The one-liner: **raw Claude Code, made warm, safe, teaching, and choreographed for a child.**

The real competitor is raw Claude Code. Out of the box it's a software-engineering assistant — dev-toned, no kid-safety, no teaching, blank-prompt, permission-prompt friction; on this machine a global "be blunt / mildly profane" CLAUDE.md would even leak onto the kid (the `kid-buddy` persona explicitly neutralizes it). Everything this product adds is the missing *kid layer*; the engine is already there.

Lean on the substrate, don't reinvent it. Claude Code already ships the heaviest items v1 wanted built: `/rewind` checkpoints + git = the "Manage & Recover" engine (v1's "build auto-checkpoint" was reinventing a built-in); SessionStart hooks = the "new powers" greeting; PreToolUse hooks = a deterministic safety gate; Output Styles = the jargon-strip. The work is *kid-skinning* these.

The genuine remaining gaps (now that Setup is scored correctly): Play & Show (4.8) + Publish (4.7) — the "show a friend" loop; Identity (4.5) — portfolio/badges; **Young Kid (3.7)** — pre-reader access (the shipped audio + a future voice-input mode); **Teen (5.7)** — a fade-to-real-code path so it doesn't feel babyish; **Educator (2.8)** — a separate product.

Workflow experience — Mid Kid (8–10)

Domain	Capability	Experience	Note
Setup & Onboarding	6	6	One-paste install + auto-settings is near-native for a CC plugin; a grown-up doing the first install is normal for any tool a 9-year-old uses — not a defect.
Make-a-Game	9	8	Complete/playable/offline fast; ceiling = coolness vs Roblox, latency.
Learn-as-you-build	8	8	The defensible wedge — no other agent teaches.
Manage & Recover	6	6	Leans on CC <code>/rewind</code> +git once <code>oops-go-back</code> is wired to it; app-DATA orphan-on-rename is the residual gap (PR #11 added guidance).
Publish & Share	7	4	Agent-driven deploy works; gated, not one-ask, no audience — the prime motivator throttled.

Ranked friction: ① no audience / not one-ask share → SHARE lever · ② blank-prompt + no first-run choreography → idea-helper/getting-started/SessionStart · ③ pre-reader access → **audio (shipped)** + voice-input · ④ `oops-go-back` not yet wired to `/rewind` → lever #1 · ⑤ unmeasurable (no telemetry) → opt-in local success signal.

Competitive position (re-anchored)

Tier	Who / role in the analysis
THE BAR raw Claude Code	Claude Code with no plugin — the default a curious family hits. Capable engine; no kid persona, no safety, no teaching, blank-prompt, permission friction. Every cell is scored against this; the product beats it on warmth, safety, teaching, choreography.
Same substrate (ports, not rivals)	OpenAI Codex CLI , Gemini CLI , Cursor (IDE), Aider — same engine class, same <i>missing kid layer</i> . Relevant as future ports (the "someday Codex" goal), not competitors.
Human substitute	A parent/older sibling who codes, in CC beside the kid — warmest+safest, doesn't scale, kid is a passenger. This product's edge: autonomy .
The road not taken	Standalone kid platforms — Scratch/ScratchJr, Code.org, Tynker, MakeCode, Roblox, Rosebud AI. What a family might use <i>instead of the whole Claude-Code path</i> . They set expectations (web instant-on, remix community, an audience) but are not the per-cell bar. Their one real lesson: kids want an audience → the gated-share lever.

First-mover. Searches of the CC plugin/marketplace ecosystem (alirezarezvani/claude-skills ~337 skills, aitmpl.com ~1367 skills, daymade/claude-code-skills) found **no kid-safe building plugin** — the one education pack is educator tooling. This is effectively **first / alone** in "kid-safe persona + teaching + build-anything plugin for Claude Code."

Honest privacy caveat: "nothing leaves the machine" is true of creations + the Brain Wall + the audio file; the kid's *conversation still goes to Anthropic's API* to generate code. Market privacy as *local-storage-of-creations* best-in-class, not *zero-network*.

Portability (the "someday Codex" goal): *most* of the asset value (skill bodies, persona prose, README pedagogy) ports with light find-replace via the shared open Agent Skills `SKILL.md` standard (Codex adopted it in late 2025); the rest is per-CLI glue (manifests, persona-as-main-thread, hooks, the `/rewind` lean). Sequence: **Codex first**, Gemini second, Aider a stretch.

Cross-cutting levers (re-ranked)

Ranked by $(\text{cells} \times \text{gain} \times \text{confidence}) / \text{effort}$. The top is cheap because the fixes are prompt-only skills or *leaning on Claude Code built-ins*.

#1 · 10.8 Kid-skin Claude Code /rewind + git (save / undo / versions)

6 cells

+2.0

High

S

Wire `oops-go-back` / "save my progress" to CC's built-in checkpoints + git, kid-skinned ("I saved a snapshot 📸"). *Lean on the built-in, don't reinvent a recovery system.* → **O3, O2.**

#2 · 10.8 Audio onboarding companion

✓ SHIPPED

6 cells

+2.0

High

S

Kid + grown-up "🎧 Listen" guides (OpenAI TTS), transcript, README links, persona backstop. Lifts the weakest in-market row (Young Kid). Ties lever #1 by formula; listed 2nd because it's already done. → **O2, pre-reader access.**

#3 · 10.5 Idea discovery + first-win + SessionStart hook

10 cells

+1.5

Med

S

`idea-helper` + `getting-started` + a "new powers" SessionStart greeting that offers 3 concrete ideas. Kills blank-prompt paralysis. → **O1, O2.**

#4 · 5.6 CSTA tagging + parent/learner progress view

8 cells

+1.0

Med

S

Map Brain-Wall concepts to CSTA; a "what they made & learned" panel. → **parent trust → O3, educator-credibility.**

#5 · 4.7 Band-D fade-to-real-code / de-baby mode

9 cells

+1.5

Med

M

A path from "buddy builds it" to "I edit real code"; a less-childish skin for teens. → **age-range-up.**

#6 · 4.2 Band-A voice-input + picture choices

8 cells

+1.5

Med

M

Voice-first prompts + picture either/or — the input complement to the shipped audio output.
→ **O2, age-range-down.**

#7 · 3.5 First-run polish + a "did it work?" check

5 cells

+1.0

Med

S

Tighten the very first session so a non-technical grown-up sees it succeed. → **O1.**

#8 · 3.5 Auto-juice + game/app breadth

10 cells

+1.0

Med

M

Screen-shake/particles/sound by default; wider genres. → **O3 (coolness).**

#9 · 2.8 One-tap grown-up-gated safe share link

8 cells

+2.0

Med

L

Pll-scrubbed play-only link near one ask (→ moderated gallery later). The audience gap. → **O3.**

#10 · 2.8 Lean on CC affordances (PreToolUse gate + Output Styles)

8 cells

+1.0

Med

M

A deterministic safety hook (hard-block publish/network/deletes) + aligning the persona with Output Styles. → **safety floor, O2.**

#11 · 2.8 Classroom edition

9 cells

+3.5

Med

XL

Rostering, dashboard, CSTA/ISTE reports, DPA, seat pricing. → **a different outcome (weekly-active-orgs)** — a deliberate market choice, not a priority.

Dropped from v1: the "no-terminal hosted path" moonshot — it contradicts the corrected frame (the terminal *is* the substrate). **Strategic (not cell-lifting):** Codex/Gemini portability — a new substrate for reach, the "someday Codex" goal.

Phased roadmap

NOW · THIS/NEXT SPRINT — LEAN-ON-SUBSTRATE + PROMPT-NATIVE

#1 wire oops-go-back to CC /rewind +git · #2  audio onboarding (shipped) · #3 idea-helper/getting-started + SessionStart greeting.

Outcome: O1/O2/O3. **Deps:** none (skills/hooks/built-ins). **Capacity:** all S — fits a sprint. **Projected:** 5.17 → ~5.6.

NEXT · THIS QUARTER — CLOSE THE REAL GAPS

#4 CSTA tagging + progress view · #6 Band-A voice-input (pairs with shipped audio) · #7 first-run polish · #8 auto-juice + breadth · #10 PreToolUse safety gate + Output Styles.

Outcome: O2/O3 + parent trust + age-range-down. **Deps:** CSTA builds on concept-cards . **Capacity:** don't co-schedule both M-safety and voice with infra. **Projected:** → ~6.2.

LATER · BEYOND — AGE RANGE, AUDIENCE, SUBSTRATE, MARKET

#5 Band-D fade-to-real-code · #9 gated share link → moderated gallery · **Codex/Gemini portability** (the "someday Codex" goal) · #11 classroom edition (different product/outcome).

Deps/capacity: gallery needs moderation infra; portability + classroom are each multi-quarter — sequence, don't parallelize. **Projected:** → ~7.0.

Target matrix (projected column averages)

Domain	Today (v2)	After Now	After Next	After Later
Setup & Onboarding	5.5	6.3	7.0	7.5
Make-a-Game	5.7	5.9	6.5	7.3
Make-an-App	5.3	5.5	6.2	7.0
Change/Remix/Fix	5.3	5.6	6.0	6.8
Play & Show	4.8	5.0	5.6	6.8
Learn-as-you-build	5.8	6.2	6.9	7.4
Identity & Progress	4.5	4.8	5.6	6.5
Publish & Share	4.5	4.7	5.4	6.6
Manage & Recover	5.0	6.6	6.8	7.2
Overall (mean of columns)	5.2	5.6	6.2	7.0

Strategic recommendations

Press the wedge: "the only AI that *builds your game/app AND teaches you while it does* — and keeps it all on your computer." First-of-its-kind on the Claude Code substrate. Lead with it; lean into **Make-an-App for young kids** (white space even beyond the substrate).

Don't reinvent the substrate. Skin /rewind +git, fire a SessionStart greeting, add a PreToolUse safety gate — free, native, and stronger than bespoke prompt-only versions.

The audience gap is the one real lesson from the standalone platforms. Kids want someone to play their thing — close it with the gated share link (then a moderated opt-in gallery), inside the existing safety model.

Two roles are different products: the Teen (needs fade-to-real-code so it isn't "cheating") and the Educator (seat pricing + rostering + DPA). Each a deliberate, separately-scoped bet.

Plan for portability so the pedagogy isn't Claude-Code-locked: keep skill bodies + persona + README tool-agnostic; isolate the CC-specific glue. Codex CLI is the natural first port.

Profiles

THE BAR

Raw Claude Code (no plugin)

Anthropic's terminal AI coding agent. It *will* build a working game/app from a plain prompt — but as a software-engineering assistant: dev tone, no kid-safety, no teaching, no first-win choreography, permission prompts. Output Styles are the documented way to strip the SWE tone; this plugin's persona does that + adds safety + teaching. [Output Styles ↗](#) · [Checkpointing/Rewind ↗](#) · [Plugins ↗](#)

SAME SUBSTRATE · PORT TARGETS

Codex CLI · Gemini CLI · Cursor · Aider

Same engine class, same missing kid layer. **Codex CLI** shares the open Agent Skills `SKILL.md` standard (cleanest port). **Gemini CLI** uses extensions + `GEMINI.md`. **Cursor** is a full IDE (more surface for a young kid). **Aider** is a dev pair-programmer (worst kid-fit). [Codex Skills ↗](#) · [Gemini extensions ↗](#) · [Open Agent Skills ↗](#)

Scratch · Code.org · Rosebud AI · Roblox

Standalone kid platforms a family might use instead of the Claude-Code path. They set expectations — web instant-on, remix culture, and an **audience/gallery** (Scratch: 164M+ projects) — but aren't the per-cell bar for a CC plugin. The audience expectation is the one worth partly meeting. [Scratch ↗](#) · [Code.org ↗](#) · [Rosebud AI ↗](#)

Audio onboarding (OpenAI TTS)

The "🎧 Listen instead" guides were generated with OpenAI `gpt-4o-mini-tts` (you own the output; ~1¢; steerable warm tone — OpenAI's terms require disclosing the voice is AI-generated, which the transcript + README links do). Higgsfield's "audio" is a separate voiceover product (reportedly ElevenLabs-backed), so ElevenLabs itself (paid) or OpenAI are the cleaner options; macOS `say` is the zero-dependency fallback. Generated once, committed as static `.mp3` — no runtime TTS, offline-first intact. [OpenAI TTS ↗](#)

Sources & references

- **Claude Code substrate (the corrected bar + what to lean on)**
- Plugins — <https://code.claude.com/docs/en/plugins> · Plugin marketplaces — <https://code.claude.com/docs/en/plugin-marketplaces> · Skills — <https://code.claude.com/docs/en/skills> · Subagents — <https://code.claude.com/docs/en/sub-agents>
- Output Styles (strip the SWE tone) — <https://code.claude.com/docs/en/output-styles> · Glossary — <https://code.claude.com/docs/en/glossary> · Hooks (SessionStart/PreToolUse) — <https://code.claude.com/docs/en/agent-sdk/hooks>
- Checkpointing / `/rewind` — <https://code.claude.com/docs/en/checkpointing> · File checkpointing — <https://code.claude.com/docs/en/agent-sdk/file-checkpointing> · How it works (snapshots before edits) — <https://code.claude.com/docs/en/how-claude-code-works> · Settings — <https://code.claude.com/docs/en/settings> · Best practices (CLAUDE.md, `/rewind`, git) — <https://code.claude.com/docs/en/best-practices>
- **Same-substrate alternatives / portability ("someday Codex")**
- Codex Agent Skills — <https://developers.openai.com/codex/skills> · Codex AGENTS.md — <https://developers.openai.com/codex/guides/agents-md> · Codex hooks — <https://developers.openai.com/codex/hooks> · Open Agent Skills standard — <https://agentskills.io/>

- Gemini CLI extensions — <https://google-gemini.github.io/gemini-cli/docs/extensions/> · Gemini CLI — <https://github.com/google-gemini/gemini-cli> · Aider — <https://aider.chat/docs/> · Cursor — <https://cursor.com/> · kid+Cursor anecdote — <https://www.haihai.ai/emma/>
- **First-mover search (no kid-safe CC plugin found)**
- alirezarezvani/claude-skills — <https://github.com/alirezarezvani/claude-skills> · aitmpl.com — <https://www.aitmpl.com/plugins/> · daymade/claude-code-skills — <https://github.com/daymade/claude-code-skills> · claude-education-skills (educator-facing) — <https://skillsllm.com/skill/claude-education-skills> · "Claude Code for Kids" classes (adjacent) — <https://codewithus.com/topic/claude-code>
- **The road not taken (expectations, not the bar)**
- Scratch — <https://scratch.mit.edu/> · stats — <https://scratch.mit.edu/statistics/> · Code.org — <https://code.org/> · App Lab — <https://code.org/en-US/tools/app-lab> · Rosebud AI (coding for kids) — <https://lab.rosebud.ai/coding-for-kids> · MakeCode Arcade — <https://arcade.makecode.com/> · Tynker — <https://www.tynker.com/> · Roblox content maturity — <https://create.roblox.com/docs/production/promotion/content-maturity>
- **Audio onboarding (feasibility, licensing)**
- OpenAI TTS (gpt-4o-mini-tts, steerable) — <https://developers.openai.com/api/docs/guides/text-to-speech> · OpenAI Terms ("you own the Output") — <https://openai.com/policies/row-terms-of-use/> · ElevenLabs output rights by plan — <https://terms.law/ai-output-rights/elevenlabs/> · Higgsfield Audio (ElevenLabs-backed) — <https://higgsfield.ai/audio> · macOS `say` → mp3 — <https://ss64.com/mac/say.html>
- **Developmental / standards / privacy (per-role)**
- NN/g Kids cognition — <https://www.nngroup.com/articles/kids-cognition/> · Teens UX — <https://www.nngroup.com/articles/usability-of-websites-for-teenagers/> · CSTA standards — <https://csteachers.org/k12standards/> · ISTE — <https://iste.org/standards/students> · FTC COPPA — <https://www.ftc.gov/business-guidance/privacy-security/childrens-privacy> · 2025 COPPA changes — <https://www.ftc.gov/news-events/news/press-releases/2025/01/ftc-finalizes-changes-childrens-privacy-rule-limiting-companies-ability-monetize-kids-data> · FERPA — <https://studentprivacy.ed.gov/ferpa> · Resnick "wide walls" — <https://mres.medium.com/designing-for-wide-walls-323bdb4e7277>
- **Platform / cost**
- Claude Code with Pro/Max — <https://support.claude.com/en/articles/11145838-use-claude-code-with-your-pro-or-max-plan> · pricing — <https://www.finout.io/blog/claude-code-pricing-2026>

v2, re-anchored to the Claude Code substrate. Working spec: `COVERAGE_MATRIX_ANALYSIS.md`. Scores are an analyst's evidence-based estimate (the product collects no telemetry by design). Audio onboarding shipped this session.